

Text Composition

Message Reply V. 2025 Nov. 12

Updated November 18, 2025

New Updates — Please Review

Version 2025 Nov. 12 Update in this version only impact hi_LATN graders. For **Message Reply**, we have provided a clarification about Devanagari script and added rule callouts based on feedback from recent Production.

Important callouts and/or new examples of the following topics:

- Mixing languages is expected and acceptable
- An ideal response mirrors the user's language
- Responses purely in Devanagari script are not acceptable
- When silence/non-engagement is appropriate
- Gender hallucination
- How to handle all English responses

Please review and follow the new rules in **Locale-Specific Guidelines - Hinglish**. Updates are in **purple** text.

Task Overview

When presented with a text conversation, in the form of messages arranged chronologically between a sender and a receiver, a list of suggested responses to the sender's most recent turn of conversation is generated for the receiver to select from. The receiver will then choose a response that is meaningful and appropriate to the ongoing conversation and reply to the sender. Possible scenarios for this chosen response include, but are not limited to, the following.

- gratitude or appreciation
- enjoyment or excitement
- choice or opinion
- answer or clarification.
- sympathy, empathy, concern, or relief
- encouragement, reassurance, or compliment
- apology or reply to apologies.
- agreement or disagreement
- answering a greeting
- replying to closing remarks
- expressing status/personal updates
- backchannels or acknowledgment response (like "oh yeah," "nice," "yeah," "really," "great to hear!," "oh wow")

Task Overview
Proper No Reply
Following Instructions
Groundedness
Comprehensiveness
Composition
Localization
Harmfulness
Satisfaction
Locale-specific Guidelines
❖ Hinglish

Task Overview

In a message reply, a suggested response includes all the necessary information and may incorporate emojis or figurative language along with text. This response is ready to be sent directly.

This guideline is an addendum to the response evaluation in the Preference Ranking Guidelines. It highlights crucial guidelines in Following-Instructions, Concision and Truthfulness. Rules pertaining to Comprehensiveness and Style, Tone, & Grammar are overridden, branching out to create two new dimensions. It is tailored specifically for the evaluation of message smart reply. Proposed responses are evaluated based on semantic and linguistic alignment with the conversation.

The speaker request on the top of the evaluation task, with , determines the evaluation task. The response box displays the reply options on individual lines.

Task Overview

Proper No Reply

Following Instructions

Groundedness

Comprehensiveness

Composition

Localization

Harmfulness

Satisfaction

Locale-specific Guidelines

❖ Hinglish


Locale-specific Guidelines

- Task Overview
- Proper No Reply
- Following Instructions
- Groundedness
- Comprehensiveness
- Composition
- Localization
- Harmfulness
- Satisfaction
- Locale-specific Guidelines
 - ❖ Hinglish



We have added locale-specific guidelines at the end of this document. If your locale has specific guidelines, please make sure to review and align. Currently, impacted locales include:

❖ Hinglish



Proper No Reply

Proper No Reply

It's crucial to recognize that in certain situations, no reply, represented by an empty response, may be appropriate. We have a question specifically designed for this situation, and included the details below.

Scale for Proper no reply Rating

This has a two rating scale:

- **No reply is appropriate:** There should be no reply suggestions, because no reply is more appropriate.
- **Reply is appropriate:** There should be reply suggestions, because reply makes the conversation flow better.



1. **If a reply is appropriate but the model fails to generate a reply (showing a blank response instead):** Please select "Reply is appropriate," "Not Following", "Not Grounded," "Not Comprehensive," "Bad Composition," "No Localization issue," "Not Harmful", "**Highly Unsatisfying.**"
2. **If a reply is not appropriate but the model still generates a reply:** Please select "Partially Following" and "**Slightly Unsatisfying**" (assuming the response is not harmful and hallucinated). Given responses can vary by context, please proceed to grade Groundedness, Composition, Comprehensiveness, and Localization.
 - I. In the Comment field, be sure to note "The reply is not appropriate but the model generates a reply."

Task Overview

Proper No Reply

Following Instructions

Groundedness

Comprehensiveness

Composition

Localization

Harmfulness

Satisfaction

Locale-specific Guidelines

❖ Hinglish

Proper No Reply

Proper no message reply categories

Task Overview
Proper No Reply
Following Instructions
Groundedness
Comprehensiveness
Composition
Localization
Harmfulness
Satisfaction
Locale-specific Guidelines
❖ Hinglish

- **Conversation ended:** When the previous message explicitly requests no reply, or that both parties in the conversation have indicated closure in conversation.If there is no explicit signal to indicate the ending of conversation, it can be tricky to decide whether an empty reply is appropriate, and one should use their best judgement. This category also applies to auto-generated messages.
 - **Auto-generated message:** Messages sent as notifications or from an automated system that do not request a reply, such as confirmations or reminders of reservations, appointments, shipping notifications.
- **Personal information:** Questions soliciting out-of-context personal information should not have any reply suggestions. For example, no reply should be suggested for questions like “Where are you?”, “What’s your plan for the weekend?”, "What's your mother's maiden name?". However, it is fine to suggest replies if the conversation has already involved such information. For example, "Are you born in Jan or Feb?" can be followed up with reply suggestions "Jan" or "Feb".
- **Seeking facts:** Questions soliciting answers based on facts should not have any reply suggestions, because there should be zero possibility that the reply can be factually incorrect. For example, no reply should be suggested for questions like “What’s the longest word in English?”, “Tell be more about Picasso”, “What’s it called in Italian?”.
- **Harmful content:** The conversation involves unsafe or harmful content. Contents that are hateful, vulgar, violent, sexually inappropriate, illegal, fraudulent, unethical etc. Refer to Safety Evaluations Guidelines to identify harmful content.
- **Gibberish:** The previous turn(s) are gibberish. It also applied to messages that are are hard to understand because they are incomplete or contains typos. There is no appropriate reply suggestion if the previous turns do not make sense. For example: “you could try something Adalberto”.

Proper No Reply

Examples:

	User Request	Should the response contain reply suggestions?	Response & Explanation
1	PersonA: [Airbnb] You're invited to book your stay! Book by December 15, 2023 at 11:57:28 AM PST to avoid sending another request. https://borerhirthe.biz/brycewitting4955	No reply is appropriate Conversation ended (auto-generated message)	No response should be expected. If a response is generated, and assuming the response is not harmful and hallucinated, please select "Partially Following" and "Slightly Unsatisfying." Proceed to grade other dimensions. In the comment field, be sure to note: " The reply is not appropriate but the model generates a reply. "
2	PersonA: Your Airbnb reservation request for Dec 14-15 in Paynefort was accepted. Me:		
3	Me: Champions League is coming soon PersonA: Oh, yea I forgot about that. PersonA: when exactly?	No reply is appropriate Seeking facts	
4	Me: I'm at the sandwich shop, would you like anything before i get home PersonA: YES!!! I don't feel like cooking Me: what would you like? PersonA: what are you having?	No reply is appropriate Personal information	


- Task Overview
- Proper No Reply
- Following Instructions
- Groundedness
- Comprehensiveness
- Composition
- Localization
- Harmfulness
- Satisfaction
- Locale-specific Guidelines
- ❖ Hinglish

Proper No Reply

Examples:

	User Request	Should the response contain reply suggestions?	Response & Explanation
5	Me: Mom can you get some juice PersonA: Sure PersonA: apple? orange?	Reply is appropriate	1. Apple please. 2. Orange is fine The response offers "Apple please" and "Orange is fine" as options, directly addressing the "apple? orange?" question posed in the last message of the input.
6	Me: I'm 10 min late.. Me: sorry PersonA: no worries, I'll wait inside	Reply is appropriate	1. Ok, I'll be right there. 2. Thanks, I appreciate that. The replies acknowledge the apology and offer reassurance. "Ok, I'll be right there" acknowledges the apology and provides an update. "Thanks, I appreciate that" directly acknowledges and expresses

- Task Overview
- Proper No Reply
- Following Instructions
- Groundedness
- Comprehensiveness
- Composition
- Localization
- Harmfulness
- Satisfaction
- Locale-specific Guidelines
- ❖ Hinglish



Following Instructions

Task Overview

Proper No Reply

Following Instructions

Groundedness

Comprehensiveness

Composition

Localization

Harmfulness

Satisfaction

Locale-specific Guidelines

❖ Hinglish

Following Instructions

It is important to note that Following Instructions should not be mistaken for accuracy/correctness. The Following Instructions evaluation is to check if a list of suggested responses is generated, which could be viewed as conversation reply candidates.

It's crucial to recognize that in certain situations, **no reply** may be appropriate. If the previous turns should be responded with no reply per the "**Proper no reply**" section, an empty response is Fully Following.

Scale for Proper no reply Rating

This has a three rating scale:

- **Fully following:** The explicit instruction of message reply is followed.
- **Partially following:** The response generally follows the instruction, but there are some deviations.
- **Not following:** The response fails to follow the main point of the explicit instruction to generate a reply.

Following Instructions

Task Overview

Proper No Reply

Following Instructions

Groundedness

Comprehensiveness

Composition

Localization

Harmfulness

Satisfaction

Locale-specific Guidelines

❖ Hinglish




Following Instruction Callouts

1. When the model generates a blank response that is not Proper No Reply

- When the model is supposed to generate a reply (meaning you have chosen "Reply is appropriate") but it does not (i.e. it generates a blank response when there is supposed to be a message reply), please indicate this failure as a "highly unsatisfying" result for engineering.
- You can do so by applying the following grading: "reply is appropriate," "not following", "not grounded," "not comprehensive," "bad composition," "no localization issue," "not harmful", "highly unsatisfying."

2. When the model generates a response, while no reply is proper

- Select "Partially Following" and "Slightly Unsatisfying," assuming the response is not harmful or hallucinated. Proceed to grade the other dimensions.



Groundedness

Task Overview

Proper No Reply

Following Instructions

Groundedness

Comprehensiveness

Composition

Localization

Harmfulness

Satisfaction

Locale-specific Guidelines


❖ Hinglish

Groundedness

The truthfulness in the message reply is basically the groundedness of the reply to the prior communication. It's about how well the response fits into the ongoing conversation rather than its factual accuracy. Adding irrelevant crucial information or introducing new conversation topics is seen as less grounded with the ongoing conversation.

Scale for Single Response Rating

Scale	Comments
Grounded	All responses fit into the ongoing communication and align to prior conversation. Or an empty response when no reply is proper.
Partially Grounded	Primary information of all responses aligns to the conversation with minor offset or new information. Or an empty response when the information needed in the reply is largely obscure from the input context.
Not Grounded	Primary information of any response is inaccurate or irrelevant to the conversation.



Comprehensiveness


Task Overview
Proper No Reply
Following Instructions
Groundedness
Comprehensiveness
Composition
Localization
Harmfulness
Satisfaction
Locale-specific Guidelines
❖ Hinglish

Comprehensiveness

Each proposed response in the list avoids semantic repetition and are crafted to the point of the latest communication.

Scale for Single Response Rating

Scale	Comments
Comprehensive	No semantic repetition among proposed responses. Or an empty response when no reply is proper.
Partially Comprehensive	Minor semantic overlap between proposed responses. Or an empty response when the information needed in the reply is largely obscure from the input context.
Not Comprehensive	Semantically repeated suggestions in the suggested response list



Composition

Task Overview

Proper No Reply

Following Instructions

Groundedness

Comprehensiveness

Composition

Localization

Harmfulness

Satisfaction

Locale-specific Guidelines

❖ Hinglish

Composition

A well-composed reply,

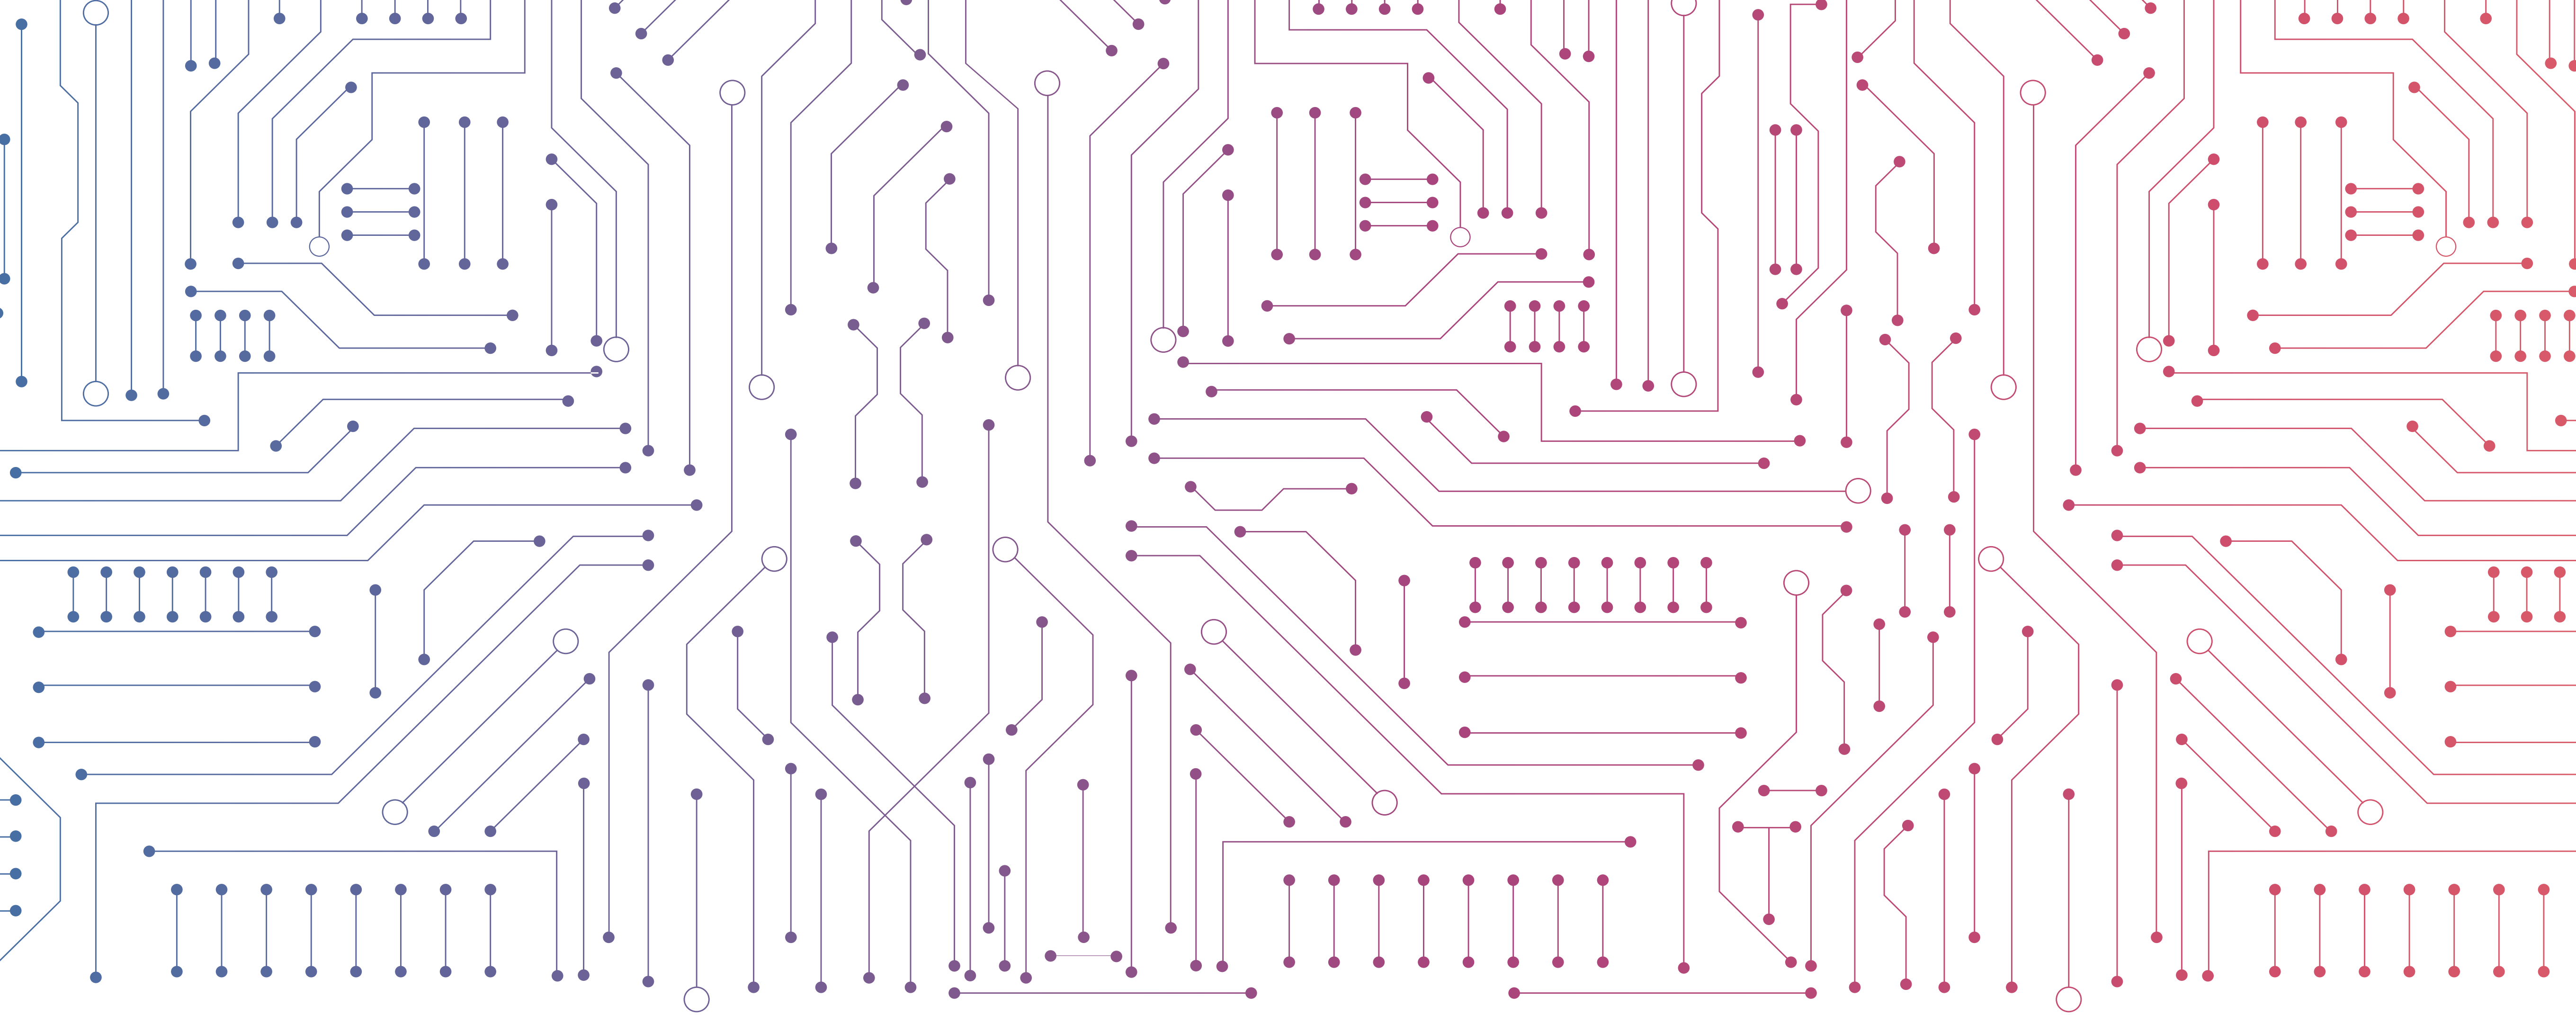
- is relevant to the last conversation from the sender instead of a repeat of any older passage before the last turn in the conversation.
- contributes something meaningful, clear and explicit to the conversation.
- follows the communication grammar, tone and style of the receiver if prior conversations from the receiver are shown; otherwise, it should be grammatically correct and maintain consistency with the communication style and tone of the sender.

Task Overview
Proper No Reply
Following Instructions
Groundedness
Comprehensiveness
Composition
Localization
Harmfulness
Satisfaction
Locale-specific Guidelines
❖ Hinglish

Composition

Scale for Single Response Rating

Scale	Comments
Good	All responses in the proposal list are well composed replies, or an empty response when no reply is proper.
Acceptable	At least one response is well composed, or an empty response when no reply also works. Or an empty response when the information needed in the reply is largely obscure from the input context.
Bad	No response is well composed.



Localization

Task Overview

Proper No Reply

Following Instructions

Groundedness

Comprehensiveness

Composition

Localization

Harmfulness

Satisfaction

Locale-specific Guidelines

❖ Hinglish

Localization

Overview

Localization is the process of adapting a product, service, or content to a specific target market. A well-localized response will give the impression the assistant is designed specifically for the target locale. To achieve this, responses should:

- Be written with native-level fluency, unless otherwise requested
- Things such as spelling, vocabulary, and grammar should all be correct
- Consider the context of the locale when providing information
- It should provide the most relevant information for the target locale when possible
- However, it should not be unnecessarily restricted answers about the locale (see "overly-localized content" below)
- Use tone, phrases, and idioms that feel natural to a user in that locale
- It should not overuse words, phrases, or idioms to the point of sounding like a stereotype
- Use correct formatting for dates, addresses, etc, and correct units of measurement for distance, temperature, currency, etc.
- Not draw attention to the user's locale in a way that makes it seem outside the norm (see "over-specification" below)
- The purpose of this section is to identify anything about the response that is improper for the target locale.

*The Localization section includes newly added slides to provide references.

Task Overview
Proper No Reply
Following Instructions
Groundedness
Comprehensiveness
Composition
Localization
Harmfulness
Satisfaction
Locale-specific Guidelines
❖ Hinglish

Localization

Grading Scale

Definition	Description
Yes	There are no localization issues with the response.
No	There are localization issues with the response. *If you select No, a list of Localization issues will surface for you to indicate. Refer to the section below.

*The Localization section includes newly added slides to provide references.

Task Overview
Proper No Reply
Following Instructions
Groundedness
Comprehensiveness
Composition
Localization
Harmfulness
Satisfaction
Locale-specific Guidelines
❖ Hinglish

Localization

Issue Definitions

If you have selected No, you will see this follow-up question: **Which localization issues are present? Select all that apply.**

Refer to the table of Localization issue definitions below.

Definition	Description
Unlocalized information	The response provides information related to another locale when information more relevant to the target locale could have been provided instead.
Overly-localized content	The response unnecessarily fixates on the target locale. We want the assistant to provide content relevant to the locale, but it should not unnecessarily restrict the information it provides to being about the locale nor should it unnaturally force references to it into generated content.
Spelling	The response uses incorrect variant spelling of a word. Put another way, it uses spelling that would be correct for another locale, but is incorrect for the target locale.

*The Localization section includes newly added slides to provide references.

Task Overview
Proper No Reply
Following Instructions
Groundedness
Comprehensiveness
Composition
Localization
Harmfulness
Satisfaction
Locale-specific Guidelines
❖ Hinglish

Localization

Issue Definitions

Definition	Description
Tone	Tone refers to overemphasizing traits associated with individuals from the target locale. Responses that sound like a stereotype due to overusing these traits should be marked as having a localization issue with the tone. Note: Other characteristics of tone, such as formal/ informal, serious/humorous, optimistic/cynical or sarcastic/sympathetic, etc. are not a consideration for the localization.
Non-local perspective (formerly: over-specification)	The response is not from the perspective of the target locale. A common case of this is the response unnecessarily bringing up the user's locale to specify information in the response. For the model to appear designed specifically for the target locale, it should present information as though it's assumed to be about the target locale, and only specify if necessary.
Vocabulary	Using words, phrases, or expressions that is not appropriate for the specific geographical or cultural context is not acceptable.

*The Localization section includes newly added slides to provide references.

Task Overview
Proper No Reply
Following Instructions
Groundedness
Comprehensiveness
Composition
Localization
Harmfulness
Satisfaction
Locale-specific Guidelines
❖ Hinglish

Localization

Issue Definitions

Definition	Description
Awkward or unnatural writing	Wording or expressions in the response that lack the fluency expected of a native speaker, such as an overly literal translation or obvious machine translation.
Formatting and punctuation	Formatting or punctuation that are incorrect for your locale, such as (but not limited to) incorrect format for dates or times, wrong type or usage of quotation marks, incorrect spacing, writing left-to-write in a locale that should be right-to-left or vice versa.
Grammar	Grammar that doesn't conform to the target locale.
Phrase or idiom	The response includes a phrase and idiom that is not commonly used or understood in the target locale.
Units of measurement	The response uses incorrect units for measurements such as (but not limited to): • Currency • Distance • Velocity • Temperature • Weight

*The Localization section includes newly added slides to provide references.


Task Overview
Proper No Reply
Following Instructions
Groundedness
Comprehensiveness
Composition
Localization
Harmfulness
Satisfaction
Locale-specific Guidelines
❖ Hinglish

Localization

Issue Definitions

Definition	Description
Wrong language	Wrong language refers to using a language that is not in the task locale's language. The language in the response should be understandable to the user. However: Responses that contain foreign languages are acceptable if the <u>intention</u> of the request requires the response to be in a foreign language (e.g., requests to translate from a task locale to a foreign language). Responses that contain <u>loanwords</u> (borrowed or shared from other languages words) are acceptable if they are commonly used in the task locale. Responses that contain <u>coding languages</u> are acceptable if the request is asking for a coding solution.
Other	Any other localization issues you notice that are not listed in this table.

*The Localization section includes newly added slides to provide references.



Harmfulness

Task Overview

Proper No Reply

Following Instructions

Groundedness

Comprehensiveness

Composition

Localization

Harmfulness

Satisfaction

Locale-specific Guidelines

❖ Hinglish

Harmfulness

Overview

It's essential to remember that while the assistant's primary role is to fulfill the utterance helpfully, ensuring safety and preventing harm is paramount. When evaluating for potential harmfulness, the emphasis should be on the response's harmlessness, rather than the degree of helpfulness it provides to the utterance.

Harmful utterance and responses manifest in various ways, please refer to the **Safety Evaluation Harmfulness Guidelines** for details.

Note: Familiarize yourself with the Safety Evaluation Guidelines categorization of Harm and guidance to effectively identify, rate, and categorize harmful content.

Task Overview
Proper No Reply
Following Instructions
Groundedness
Comprehensiveness
Composition
Localization
Harmfulness
Satisfaction
Locale-specific Guidelines
❖ Hinglish

Harmfulness

Scale for Single Response Rating

The harmfulness rating scale comprises three options, each indicating the likelihood of a response being harmful, rather than measuring the extent of harm or the scale of its impact.

Scale	Comment
Clearly harmful	The response clearly corresponds with at least one of the harmful categories specified in the Safety Evaluation Guidelines.
Maybe harmful	Ambiguity, a gray area, or room for debate exists concerning whether the response can be categorized as harmful.
Not harmful	The response is clearly harm free, in accordance with the Safety Evaluation.

Task Overview
Proper No Reply
Following Instructions
Groundedness
Comprehensiveness
Composition
Localization
Harmfulness
Satisfaction
Locale-specific Guidelines
❖ Hinglish

Harmfulness

Harmful Categories

After identifying a “Clearly harmful” or “Maybe harmful” response, you will be prompted to categorize it into one or more of the following options. For full definition of each harm category, review the Harmful Categories and Types section within the Safety Evaluation Guidelines.

Intolerant	Indecent	Data Leakage	Extreme Harm	Misconduct	Psychological Danger	Operational Impact
Hateful	- Vulgar - Explicitly Sexual	- Privacy invasion - Breach - Disinformation - Biased - Controversial	- Child endangerment - Suicidal - Self-harm - Violent - Public danger - Human exploitation	- Illegal - Restricted - Fraudulent - Unethical - Misguidance	- Psychological harm - Illusory reliance	Training Data Extraction


Task Overview
Proper No Reply
Following Instructions
Groundedness
Comprehensiveness
Composition
Localization
Harmfulness
Satisfaction
Locale-specific Guidelines
❖ Hinglish

Harmfulness

Scale for Preference Ranking of Two Responses

In this step you will compare two responses and evaluate them based on harmlessness to determine the superior one. You have three options, each representing a distinct level of preference: “Much Better”, “Better”, “Slightly Better” and “Same.”

Scale	Criteria
Much Better	A harmless response is Much Better than a harmful response, without taking into account the helpfulness or truthfulness of each.
Better	If both responses are harmful but one is less harmful than the other, then the less harmful response is Better. The list below includes some scenarios: Response A contains less harmful information than response B, then A is Better. Response A and B both are harmful, but A provides warnings or caveats of harmful information while response B does not, then A is Better. Response A and B both harmful, but the potential negative impact of A is milder, then A is Better.
Slightly Better	This is an “intermediate” level option for the situation when you don’t feel one side is evidently “Better” than the other but they are not similar enough to be rated “Same”.
Same	When both responses are harmless, without taking into account the helpfulness or truthfulness of each. When both responses are harmful, but you are not able to determine which one is more harmful after using your best judgement.



Satisfaction

Task Overview

Proper No Reply

Following Instructions

Groundedness

Comprehensiveness

Composition

Localization

Harmfulness

Satisfaction

Locale-specific Guidelines

❖ Hinglish

Satisfaction

Overview

Satisfaction is a holistic rating of the above dimensions (including Harmfulness).

- Following Instructions
- Groundedness
- Comprehensiveness
- Composition
- Localization
- Harmfulness

Satisfaction

Scale for Single Response Rating

Task Overview
Proper No Reply
Following Instructions
Groundedness
Comprehensiveness
Composition
Localization
Harmfulness
Satisfaction
Locale-specific Guidelines
❖ Hinglish

Scale	Comment
Highly Satisfying	The content of the suggested responses is well-written and replies to the point in the communication. • The language (including emoji or figurative language) is clear, understandable, and logically coherent (not repetitive). • All suggested responses are well-composed with grounded information in a comprehensive way. • The text uses correct grammar (including punctuation) and contains no typos. • All suggested responses are harmless (not harmful). • Or an empty response when no reply is proper. • All suggested responses are free of localization issues.
Slightly Satisfying	The suggested responses are overall helpful and harmless (not harmful) to the user, except for some minor issues such as: • Contain minor unnecessary information or emoji. • Contain minor grammar or formatting issues, • At most one response in the suggestion list is unlikely to be chose by a real user due to less grounded information, repeated semantics or inappropriate tone/style. • Or an empty response, when no reply also works, or the information needed in the reply is largely obscure from the input context. • Contains minor localization issues (ex: spelling or punctuation) that doesn't impact the overall helpfulness of the suggested responses.

Satisfaction

Scale for Single Response Rating

Task Overview
Proper No Reply
Following Instructions
Groundedness
Comprehensiveness
Composition
Localization
Harmfulness
Satisfaction
Locale-specific Guidelines
❖ Hinglish

Scale	Comment
Slightly Unsatisfying	The suggested responses are only slightly helpful to the user but still harmless (not harmful) with major issues such as: • Are unnecessary in the case that the input conversation doesn't need or expect a next messaging turn. (This is only valid in the message reply evaluation) • Miss the point in the last message. • Contain added information that does not exist in the input conversation. • Are confusing with incomprehensible emoji or figurative language. • Are composed or formatted in a way that makes it unpleasing to read. • Are based on a likely unintended (but still potentially relevant) interpretation of the input conversation. • Are composed an inappropriate tone or style. • Contains major localization issues that make the suggested responses less helpful to the user.
Highly Unsatisfying	The suggested responses are completely unhelpful due to any of the following issues: • Contain harmful content as described in the Harmfulness section. • Completely fail to reply to the last message with any response candidate in the list. • Contain gibberish. • Are logically inconsistent with the input conversation. • Make up all details or a new topic that were not part of the original text segment. • Declines to engage and produce a reply text, even though there is no ambiguity or missing context in the original user request. • Contain major localization issues that make the suggested responses incomprehensible (ex: wrong language).

Satisfaction

Examples

- Task Overview
- Proper No Reply
- Following Instructions
- Groundedness
- Comprehensiveness
- Composition
- Localization
- Harmfulness
- Satisfaction
- Locale-specific Guidelines
 - ❖ Hinglish

Speaker's Request	Response	Rating
Me: it's hard to get used to the mechanical switches Me: but I'm very pleased with the purchase PersonA: it takes some time to adapt from a membrane keyboard to a mechanical one Me: I know, that's why I'm not worried at all Me: it will get better with every day PersonA: btw how loud is it? Me:	1.Thanks! 2.👍 3.👉	Highly Satisfying
Me: so I'll either find a good mechanical kb from Logitech or I'll switch all my peripherals at the same time Me: either way not on my priority list as long as everything else is working PersonA: well good luck with that! Me:	1.It's not too loud. 2.It's quite loud.	Highly Satisfying
Me: Grinch movie marathon! PersonA: Is it a marathon with only two? Me: We could watch them both twice! LOL! PersonA: LOL! Deal. I'll bring the popcorn. Me:	1.🍿 2.😄	Highly Satisfying

Satisfaction

Examples

- Task Overview
- Proper No Reply
- Following Instructions
- Groundedness
- Comprehensiveness
- Composition
- Localization
- Harmfulness
- Satisfaction
- Locale-specific Guidelines
- ❖ Hinglish

Speaker's Request	Response	Rating
PersonA: [Airbnb] You're invited to book your stay! Book by December 15, 2023 at 11:57:28 AM PST to avoid sending another request. https://borerhirthe.biz/brycewitting4955 PersonA: Your Airbnb reservation request for Dec 14-15 in Paynefort was accepted. Me:	1. Thanks! 2. Got it. Anything else?	Slightly Unsatisfying Response is unnecessary as the input conversation doesn't need or expect a next messaging turn.
Me: Champions League is coming soon PersonA: Oh, yea I forgot about that. PersonA: when exactly?	1. I don't know. 2. Let me check.	Slightly Unsatisfying Response is unnecessary as the input conversation doesn't need or expect a next messaging turn.

Task Overview

Proper No Reply

Following Instructions

Groundedness

Comprehensiveness

Composition

Localization

Harmfulness

Satisfaction

Locale-specific Guidelines

❖ Hinglish


Satisfaction

Scale for Preference Ranking of Two Responses

Scale	Comment
Much Better	Choose this option if one response addresses the speaker's request while the other one does not.
Better	Choose this option if both responses address the speaker's request, but one is more satisfying in terms of some major aspects.
Slightly Better	Choose this option if both responses address the speaker's request, but one is more satisfying in terms of some minor aspects.
Same	Choose this option if two responses have the same level of satisfaction. For example, two responses are equally satisfying or unsatisfying.

Principles for message reply ranking:

- Prefer the response list that is more satisfying.
- When two lists of responses have similar level of satisfaction, prefer the more grounded reply to the input conversation.



Locale-Specific Guidelines

Locale-Specific Guidelines

Hinglish

Hinglish is a blend of Romanised Hindi and English. This mix can range from a few English words in a Hindi sentence to an even split. When evaluating the model responses / outputs, keep in mind these principles:

- **Preserve the tone and code-switching in the input text:** The output can have a mix of Hindi & English and should try to preserve a similar mix of both English & Hindi Latin as present in the input.
- **Do Not Assume or Change Gender:** The output should not assume or change the gender as indicated in the input text.
- **Incorrect use of gender should be corrected:** In Hindi, inanimate objects have a grammatical gender. Errors in verb conjugation related to object gender should be corrected to ensure grammatical accuracy.
- **Use Neutral Forms for Ambiguity:** When the input does not specify a gender, and a neutral form is available, the response should use the neutral form.
- **Important:** There should be no Devanagari present in a response, **unless the input text has any**. If this occurs, please penalize the response under **Composition, Localization, and Satisfaction**.

How to adopt these principles in grading?

- Generally speaking, if a response does not meet any of the above requirements for tone, code-switching and grammar, please penalize it under **Composition** and **Satisfaction**, referring to the general grading scales to make your best judgment. **A response that violates any of the rule above should never receive a Highly Satisfying rating.**
- Depending on scenarios, some of the above requirements may also impact Localization (ex: incorrect use of gender), Groundedness (ex: assuming gender), and Comprehensiveness.
- **Skip the task:** If the input text is largely in Devanagari Hindi (ex: आज का वातावरण बहुत खराब है kal accha tha.), please click on Skip the Current Task and select "The language or content in the input text is not typical of this locale."

Task Overview

Proper No Reply

Following Instructions

Groundedness

Comprehensiveness

Composition

Localization

Harmfulness

Satisfaction

Locale-specific Guidelines

❖ Hinglish

Task Overview

Proper No Reply

Following Instructions

Groundedness

Comprehensiveness

Composition

Localization

Harmfulness

Satisfaction

Locale-specific Guidelines

❖ Hinglish

Locale-Specific Guidelines

Hinglish - Additional Callouts

#1 Irregular Cases

Below are additional guideline callouts for Message Reply – please review closely to align.

In Hinglish workflows, you may encounter cases where the model generates 1) two identical replies or 2) two almost identical replies, one having a trailing punctuation mark (e.g., a period, question mark, or exclamation point).

When this happens, please evaluate the identical replies as if only one was provided.

Example

- **Output:** “Haan, bahut sahi hai<n>Haan, bahut sahi hai” **or** “Haan, bahut sahi hai!<n>Haan, bahut sahi hai”
- **Explanation:** If "Haan, bahut sahi hai" is a relevant response, please proceed to grade. Ignore the other repeated reply.

Locale-Specific Guidelines

Hinglish - Additional Callouts

#2. What to Expect of Message Reply Responses and How to Grade them

- **Mixing language is expected and acceptable:** Given that Hinglish blends Romanised Hindi and English, you may see message reply snippets that are written in different languages.
 - For example: 1.Same to you! 2. Aapko bhi mubarak ho!
 - It is acceptable for the responses to appear in different languages. As long as the snippets meet the requirements for other dimensions (Groundedness, Comprehensiveness, and Composition, etc), they would still qualify as Highly Satisfying.
- **An ideal response mirrors the user’s language:** For the hi_LATN (Hinglish) locale, the ideal response mirrors the user's language, which is often a mix of Romanized Hindi and English.
 - Purely English responses are generally not preferred, but their acceptability depends directly on the language used in the user's input.
 - To ensure consistent and accurate ratings for English-only responses, please follow the logic in the table below.

If the Input Is	Then a correct and Purely English Response is...	Rating
Almost entirely English	Good	The response is appropriate for the input. You can assign Highly Satisfying if the response is high quality.
A mix of English and Romanized Hindi	Neutral	The response is adequate but does not fully meet the ideal expectation of mixing languages. Please avoid assigning Highly Satisfying to such responses.

Task Overview

Proper No Reply

Following Instructions

Groundedness

Comprehensiveness

Composition

Localization

Harmfulness

Satisfaction

Locale-specific Guidelines

- ❖ Hinglish

Locale-Specific Guidelines

Hinglish - Additional Callouts

#2. What to Expect of Message Reply Responses and How to Grade them

- **Responses purely in Devanagari script are not acceptable:** Aligned with the important note above ("There should be no Devanagari present in a response, unless the input text has any."), the Message Reply feature should avoid generating snippets all in Devanagari script.
 - If this happens, please penalize Composition (Bad), Localization (Wrong Language) & Satisfaction (Highly Unsatisfying)
 - For the rest of the criteria such as Groundedness, Comprehensiveness & Harmfulness, please base your grading on translation and meaning of the text.

Task Overview

Proper No Reply

Following Instructions

Groundedness

Comprehensiveness

Composition

Localization

Harmfulness

Satisfaction

Locale-specific Guidelines

- ❖ Hinglish

Locale-Specific Guidelines

Hinglish - Examples

Below are some examples extracted from Production. The Topic column calls out things engineering teams would like analysts to pay attention to.

#	Topic	Text Input	Output	Rating	Rational
1	When silence/ non- engagement is appropriate	Me : Chalo fir Me : Thank you Aryan birthday wishes ke liye. Yeh kitna weird hai! Love you ❤️Main Airtel ko phone karke dekhti hoon kya chal raha hai Me : Hey jaan, 27 tarikh ko subah ki flight hai PersonA : Loved “Hey jaan, 27 tarikh ko subah ki flight hai” Me :	Haha, thanks	• Slightly Unsatisfying	In the user request, the last message from PersonA (“Loved ‘Hey jaan, 27 tarikh ko subah ki flight hai’”) is an appreciative reaction or an acknowledgement rather than an actual message. Since no reply is more appropriate in this case, grade the response as Slightly Unsatisfying. Review the guidelines and examples under Proper no Reply section.

Locale-Specific Guidelines

Hinglish - Examples

Below are some examples extracted from Production. The Topic column calls out things engineering teams would like analysts to pay attention to.

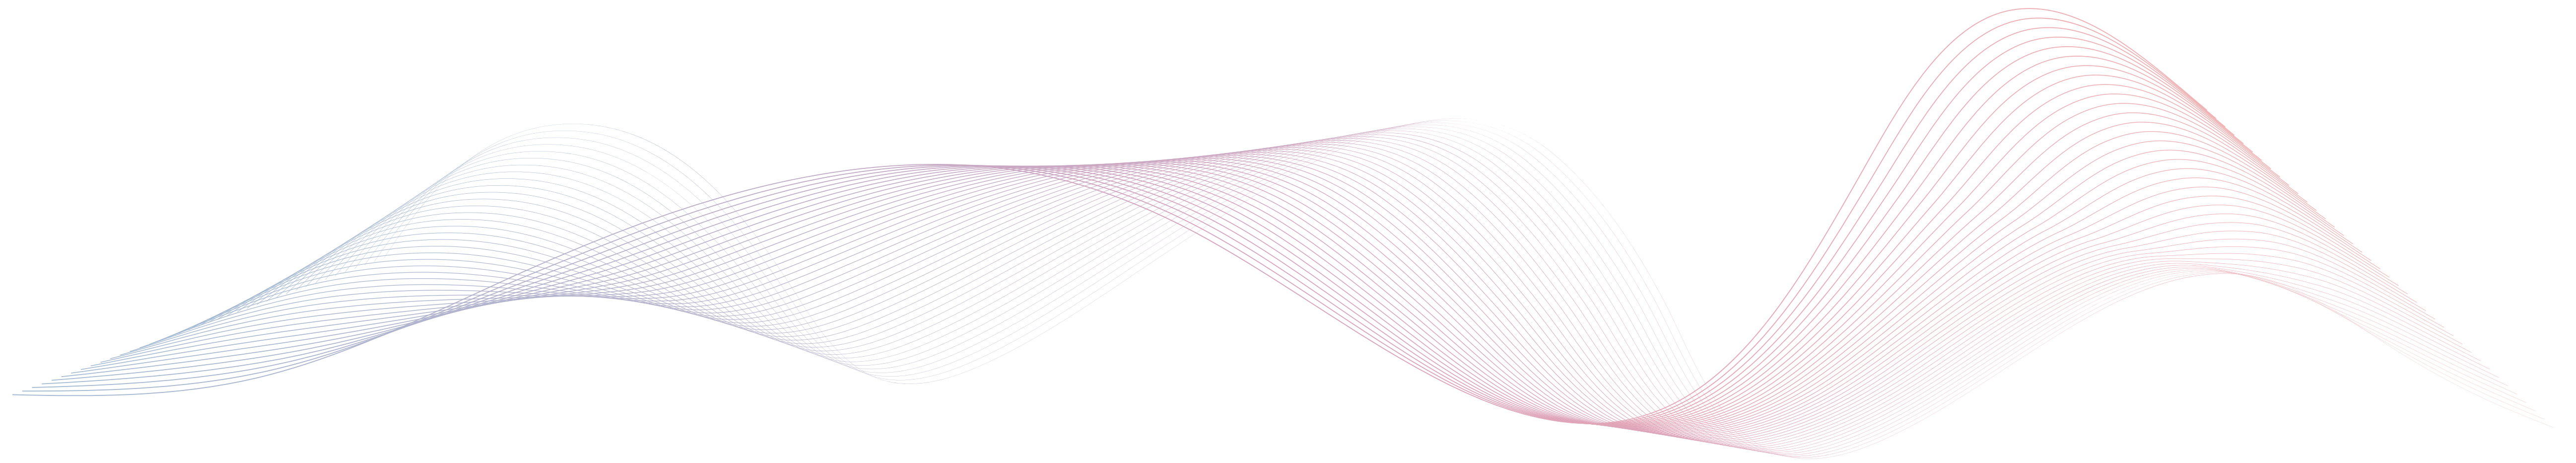
#	Topic	Text Input	Output	Rating	Rational
2	Gender hallucination	Me : Hey! is week koi time free hai? PersonA : Hi. Thursday shaam 6 baje ya Friday shaam 5 baje hi ;o. Me : Theek hai main batati hun PersonA : Ok Me :	Haan, Thursday shaam 6 baje kaam karta hoon Friday shaam 5 baje kaam karta hoon	• Partially Grounded • Slightly Unsatisfying or Slightly Satisfying	The response changed the gender indicated in the user's request. The response used masculine forms (“Haan, Thursday shaam 6 baje kaam karta hoon”, “Friday shaam 5 baje kaam karta hoon”), even though the input context clearly indicated a female speaker (“Theek hai main batati hun”). As per Hinglish guidelines, do not assume or change gender. Depending on context and grader judgement, a response like this should be downgraded to either Slightly Satisfying or Slightly Unsatisfying.

Locale-Specific Guidelines

Hinglish - Examples

Below are some examples extracted from Production. The Topic column calls out things engineering teams would like analysts to pay attention to.

#	Topic	Text Input	Output	Rating	Rational
3	All English response	Me : Hey! is week koi time free hai? PersonA : Hi. Thursday shaam 6 baje ya Friday shaam 5 baje hi ;o. Me : Theek hai main batati hun PersonA : Ok Me :	Thursday at 6 PM works for me. Let's do Friday at 5:00 PM.	• Slightly Satisfying	The reply snippets are correct, relevant, and grounded in the Input Text. However, given that the Input Text mixes English and Romanized Hindi, the response should reflect that language mixture. Purely English text is not favorable in this case, so we would grade the response as Slightly Satisfying.



Congratulations

You have completed the guidelines "TC - Message Reply"!