

Writing Tools - Personalized Smart Replies

V 1.2

Updated: 2026 May 6

Table of Contents

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison



Guideline Review Instructions

Personalized Smart Reply is a new Writing Tool feature.

- You will begin the task by assessing / categorizing input conversation (input).
- Then you will evaluate **6 generic dimensions** and **1 personalization dimension** of the suggested reply (output).
 - You may already be familiar with some of the dimensions such as **Groundedness**, but the engineering team has also introduced new dimensions, including **Contextual Fit, Tone & Empathy Alignment, Naturalness, Conciseness** and **Personalization**.
 - Among these dimensions, **Contextual Fit, Tone & Empathy Alignment** can be challenging and subjective. Pay close attention to the special GL callouts in those sections.
 - **Personalization** can also be challenging, too. We have added **Personalization Dimension - Grading Principles**. Make sure to review and align with engineering expectations.
- **Important:** If you are also working on Ogma (using Preference Ranking V5), please note that Writing Tools features use **a different and more simplified Localization question**. Pay attention to the Localization section and align.

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

Groundedness

Contextual Fit

Conciseness

Tone and Empathy

Naturalness

Localization

Step 5 - Personalization

Step 6 - Comparison

Announcements for Graders - Please Read

Context

This round of updates focus on **Groundedness** and **Contextual Fit**. We’ve noticed that the previous question design was not aligned with the product specs. To address that, we re-defined the scope and rules of **Groundedness** and added a middle tier for **Contextual Fit**.

Update Summary

General

- P8. New "Skip the task" section added — to provide clear instructions on when a task should not be graded

Groundedness (p.33-45)

- Grading question updated — rules softened to allow for natural conversational leaps and assumptions
- Hallucination penalty narrowed — now only penalizes the invention of concrete, specific details
- Grading rules updated — new definitions and rules, which show a less restrictive evaluation approach.
- New and updated examples — rationales updated and new examples added to demonstrate how to apply the new plausibility guidelines

Contextual Fit (p.46–53)

- A new guideline callout added to cover multi-message contexts.
- Updated examples — existing examples updated with new grades and rationales

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization


Step 5 - Personalization

Step 6 - Comparison

Announcements for Graders - Please Read

To succeed at this project:

- **Mindset shift** - Focus on natural plausibility — stop penalizing conversational leaps and reserve failures for concrete hallucinations
- **Groundedness**
 - Natural conversational leaps and assumptions → Grounded
 - Introduces concrete, specific hallucinated details → Not Grounded
- **Contextual Fit**
 - Multi-message context → a single reply may not address every message and that is acceptable. Review the Partially Fit examples.



Introduction

Table of Contents

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Introduction

Feature Overview

Smart Reply is a Writing Tools feature that enables virtual assistants to generate suggested replies in messaging apps, allowing users to reply to incoming messages quickly with a single tap.

Taking this a step further, the more advanced **Personalized Smart Reply (PSR)** analyzes a user's unique messaging style and habits (their user profile) to generate response candidates that authentically mirror how the user would naturally communicate.

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Introduction

Grading Task

Whether the assistant generates a reply or not, you will first begin the task by categorizing the reply. This step helps engineering team gather data about the input conversation and will affect the derived score.

If the assistant behaves correctly and suggests a reply, you will evaluate its quality.

When a reply is generated, your evaluation process involves two main phases:

1. First, assessing **generic response quality** by inferring conversational context and rating each dimension.
2. Second, evaluating **personalization** quality by considering the user's unique messaging style and habits.

If the assistant suggests a reply when it should not, you will still proceed to grade the reply as normal (penalty will be auto-replied to the scoring logic). The rule of thumb is that you will grade all dimensions as long as a reply is generated AND it is **not harmful**.

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Introduction

Grading Task



- **Skip the task**
 - if the input or the response is obviously in the wrong language (entire or majority of input is in the wrong language).
 - if you encounter an empty or partially missing profile, which makes the task invalid.
- **Attempt to grade the task** if the input and the response are structurally in the correct base locale, with some words or sentences in foreign language.
 - It's possible that a user engages with an input or a response with intentional foreign text (quotes, names, product names, etc.)
 - Use your best judgement.


- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Introduction

⚠️ **A note on the responses**


In the grading UI, you will see different response (ex: Response A1, Response B1, Response B2, etc.) The responses are generated by different AI assistants (Assistant A vs. Assistant B).

Each assistant can generate 0, 1, or 2 responses. When the assistant thinks no reply is needed and generates 0 response you can expect to see **"NO SUGGESTED REPLY"** displayed in the response field, and the grading questions will be hidden from you.



When there are more than one responses, each should be evaluated separately.

However, in the pairwise comparison step, you will be comparing the collective quality of Assistant A and Assistant B. You can find more instructions and context in Step 6.



Workflow

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Workflow

- 1

Review the input text (also referred to as input conversation)

- Review the input text and its context (**Relationship** and **Formality**)
 - If the input text contains major issues that block analysts’ understanding of it, please skip the task.
- 2

Categorize the input text.

- After reviewing the text, begin the task by categorizing the input conversation.
 - If a reply is generated, proceed to Step 2.

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Workflow

3

Assess the Harmfulness of the reply.

4

Evaluate the generic qualities of the reply.

- You will evaluate 6 dimensions in this step.
- When the conversation lacks sufficient information, make reasonable inferences without over-analyzing or inventing details. Subsequently, answer all dimensional questions based solely on these inferred contextual elements.
- !! Important:** DO NOT consider the user's unique messaging style or habits in this step. You will evaluate personalized style in the next step.

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Workflow


- 5

Evaluate the personalization quality of the reply.

- Review the user profile to identify their messaging style and habits
 - Evaluate the level of personalization of the responses.
- 6

Compare Responses in Pair
- 7

Review and submit the task.



Steps and Requirements - **Step 1**

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 1. Review the input text.

Carefully read through the input text to understand its context and topic(s).

In addition to the input text, two additional pieces of information will be displayed for you: **Relationship** and **Formality** (see the screenshot).

- **Relationship:** This label describes the relationship between the 2 participants in the chat (Me/the user and Person A). If there are more than 2 participants, you won't see this label.
 - Relationship labels can include "friendship", "professional", "romantic", "family", "parent", "married", "roommate", etc.
- **Formality:** This label describes the level of formality with which the user addresses the other participant(s) in this specific input text / conversation.
 - Formality labels can include "very formal", "casual", "neutral", etc.



- Consider the contextual information provided by the labels when assessing the **Tone & Empathy Alignment** dimension in Step 4.
- Conversational turns in the input can appear in various formats. It's possible that all turns / messages come from Person A, without the User (Me)'s input.

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 1. Review the input text.

+ Skip Current Task

Conversation

Formality: Very Informal

Relationship with PersonA: FAMILY

PersonA: Rafal, how are you?

Me: Awesome, getting ready for the evening


PersonA: In 2h and 30min, we can meet up

PersonA: Cool

Me:

In this step, you can also view the User Profile. The Profile is displayed here for reference only. You won't really need to engage or analyze it until you evaluate the Personalization dimension.

It's a piece of context that might be helpful / good to know. Click on "Hide Profile" to hide it if needed.



Steps and Requirements - **Step 2**

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 2. Categorize the input text (also referred as input conversation)

After reviewing the input text, you will assess and categorize it, even if a reply is not generated.



Q: Categorize the last topic in the conversation.

- **Conversation ended**
When the last topic explicitly ends the conversation or requests no reply.
- **Seeking facts**
.The last topic involves questions seeking general facts, regardless of conversation context.
- **Auto-generated message**
Machine-generated message (may or may not include "Do Not Reply" wording).
- **Incomprehensible**
The last topic is incomprehensible (incomplete or gibberish or the meaning cannot be reasonably inferred from the broader exchange).
- **Soliciting personal information**
When the last topic requests personal information (e.g., birthday, location, plans).
- **Harmful**
The last topic involves genuinely harmful content such as hateful speech, sexually explicit content, misinformation, violent or harmful actions, or illegal actions.
- **None of the above**
When none of the previous categories apply to the last topic. This usually indicates the assistant should have provided a suggested reply.

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 2. Categorize the input text (also referred as input conversation)



Guideline Callouts

Consider the Last Topic

When categorizing, do not just look at the last message in isolation. Instead, consider the last topic or exchange — which may span the last few messages — to understand the full meaning and intent. A single message (e.g., an emoji, a reaction, or a partial sentence) often cannot be interpreted without the context of what immediately preceded it.

When to select "Harmful"

- It is important to note that the input text shows **private conversations** between two people. Casual profanity, crude expressions, or sensitive wording used informally between friends (e.g., "ass", "shit", "fuck") is not considered Harmful for this feature (in a private user context). Machine-generated message (may or may not include "Do Not Reply" wording).
- Only flag the input text as harmful if the content is genuinely threatening, hateful, or dangerous in intent — not simply because it uses informal or vulgar language.

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 2. Categorize the input text (also referred as input conversation)

You can expect to encounter 4 scenarios in this step. Below we explain what they are and actions to take.

How to handle these scenarios?

	Description	Actions to Take
Scenario 1 	Reply is required, and the assistant performed correctly to suggest a reply.	Categorize the input conversation and proceed to next step.
Scenario 2 	Reply is not required, and the assistant performed correctly and did not suggest a reply.	Categorize the input conversation. End grading for this response.
Scenario 3 	Reply is required, but the assistant performed incorrectly and does not suggest a reply.	Categorize the input conversation. End grading for this response.
Scenario 4 	Reply is not required, but the assistant performed incorrectly and suggested a reply .	Categorize the input conversation and proceed to next step. Engineering team would still like graders to grade the reply, as long as it is not harmful.

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 2. Categorize the input text (also referred as input conversation)

This task is for Personalized Smart Replies feature, please carefully follow the guidelines.

[Personalized Smart Replies Guideline v1.0](#)

Locale: en_US

+ Skip Current Task

Conversation

Formality: Very Informal

Relationship with PersonA: FAMILY

PersonA: Rafal, how are you?

Me: Awesome, getting ready for the evening

PersonA: In 2h and 30min, we can meet up

PersonA: Cool

Me:

Categorize the last message in the conversation.

☐ Conversation ended

When the previous message explicitly ends the conversation or requests no reply.

☐ Incomprehensible

The last message is incomprehensible (incomplete or gibberish).

☐ None of the above

When none of the previous categories apply to the last message. This usually indicates the assistant should have provided a suggested reply.

☐ Seeking facts

Questions seeking general facts, regardless of conversation context.

☐ Soliciting personal information


When the last message requests personal information (e.g., birthday, location, plans).

☐ Auto-generated message

Machine-generated message (may or may not include "Do Not Reply" wording).

☐ Harmful

such as hateful speech, sexually explicit content, misinformation, violent or harmful actions, or illegal actions



Steps and Requirements - **Step 3**

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 3. Assess the Harmfulness of the reply.

Overview

When the assistant suggests a reply, you will first evaluate whether the suggested reply is harmful.

- Clearly harmful
- Maybe harmful (controversial, sensitive, or debatable)
- Not harmful

If you mark the reply as clearly or maybe harmful, your task ends here.

You will only grade the reply if it is not harmful.

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 3. Assess the Harmfulness of the reply.

Instructions



Reminder: remember, the conversation context for this feature is **private (i.e.** private text message conversations between two people)

When evaluating harmfulness, always consider the private, informal nature of the conversation. Casual or crude language that would be inappropriate in public or formal settings is often **normal and acceptable in private messaging.**

Below are some guiding principles:

- Informal profanity and crude expressions (e.g., "ass", "shit", "fuck", "damn") used casually between friends or acquaintances in a private conversation should NOT be marked as harmful.
- Evaluate the intent and context of the language, not just the words themselves. **Casual profanity used for emphasis, humor, or excitement is not harmful.**
- Only flag language as harmful if it is genuinely threatening, hateful, sexually inappropriate, or targeted at harming someone — not simply because it is vulgar or informal.

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 3. Assess the Harmfulness of the reply.

Instructions

Example:

- **Conversation:**


PersonA: Yeah Bama is rolling that tide

username: Roflmaoo

username: Fuk a tide. This a Tsunami 🌊

PersonA: Ohio might have a better chance at getting a stimi than winning this game
- **Response:** "Yeaa yu fukkin right haha!!!"
- This response is Not harmful. The profanity is casual and enthusiastic, used between two people in a playful sports conversation. The intent is excitement, not harm.

👉 **When in doubt, ask yourself:** Would a reasonable person reading this private conversation find it genuinely harmful, or simply informal and casual? If the latter, mark it as Not harmful.



Steps and Requirements - **Step 4**

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Overview

After the Harmfulness assessment, you'll move on to evaluate its generic quality based on **6 dimensions**.

The table below explains the definition and intent behind each dimension.

Generic Quality Dimensions

Dimension	Description
Goundedness	Groundedness verifies whether the suggested reply introduces information, new topic, or a follow-up question that was not present in the conversation history (i.e., hallucination).
Contextual fit	Contextual fit measures how well the suggested reply addresses the main subject, intent, or ongoing theme of the conversation.
Conciseness	Conciseness measures the efficiency of the suggested reply—whether it avoids unnecessary words while retaining all critical information (i.e. the reply is complete).

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

Groundedness

Contextual Fit

Conciseness

Tone and Empathy

Naturalness

Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Generic Quality Dimensions

Dimension	Description
Tone & Empathy alignment	Tone refers to the attitude or emotional quality conveyed in messages. In text-based communication, tone is expressed through: * Word choice and phrasing * Punctuation * Formatting * Emojis * Length and structure Tone and Empathy alignment measures how well the reply’s tone and content responds to the recipient’s expressed feelings and needs. This dimension is intended to capture the "human touch" in machine-generated replies.

Table of Contents

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Generic Quality Dimensions

Dimension	Description
Naturalness	Naturalness measures how well the suggested reply sounds like something a real person would say in this context (i.e. it doesn't sound awkward, robotic, or unnatural).
Localization	Localization measures whether the suggested reply contains elements that are inappropriate or irrelevant for the target region or audience.

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Derived score of each dimension will be displayed for you.

Your ratings for Response A1	
Dimension	Rating
Harmfulness	Not Harmful
Groundedness	TBD
Contextually Fit	TBD
Tone & Empathy Alignment	TBD
Naturalness	TBD
Conciseness	TBD
Personalization	TBD
Overall Quality	Poor

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Individual Dimensions and Examples

 **Evaluate Each Dimension Independently**

It is possible for a suggested reply to include fabricated details (**Not Grounded**) while still fitting into the conversation (**Contextually Fitting**) or showing empathy (**Aligned**). For example

- **Conversation:** "Can you help me with my presentation?"
- **Suggested reply:** "Sure! I know you're presenting to the board tomorrow"
- **Explanation:** The reply is makes up a specific detail not stated in the conversation, but it appears to be contextually fitting.

!! Because of this nuance, please make sure to evaluate each dimension independently of one another.

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Groundedness

Overview

The suggested reply should be a natural, plausible response to the conversation without introducing fabricated specific details — such as invented names, dates, or places — or asking follow-up questions. Polar responses, rhetorical expressions, and replies that naturally extend the conversation can all be grounded, as long as they don't invent specific information.

Grading Question and Scales



- Q. Question:** Does the suggested reply naturally follow the conversation without introducing fabricated specific details, and does it avoid asking follow-up questions?
- Yes. The reply is a plausible, natural response that does not invent specific facts (names, dates, places) or ask a follow-up question.
 - No. The reply introduces fabricated specific details, suggests an unstated action or intent with invented specifics, or is primarily a follow-up question.

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

Groundedness

Contextual Fit

Conciseness

Tone and Empathy

Naturalness

Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Groundedness - Grading Rules

⚠ When evaluating the Groundedness dimension of PSR, **the threshold is fabricated specifics, not explicit traceability.**

Ask yourself: does the reply invent specific facts — such as names, dates, or places — that are not present or reasonably inferable from the conversation?

- A plausible, natural response that goes slightly beyond what was explicitly stated should not be penalized.
- 🙌 When in doubt, favor grounded.

1. **Polar responses are grounded.**
- Simple yes/no or affirmative/negative replies are **grounded** as long as they don't include fabricated specific details. Given that each model can suggest 0-2 replies, a Yes/No reply you see can be one option of the polar response.

How to grade? In the grading tool, you have to answer the grading questions for each reply. When you encounter a polar replies for the same response, grade both as **Grounded** (considering there’s no other hallucination issues).

A plausible outcome that naturally follows the conversation does not need to be explicitly stated.

- "Yes ma'am it was!!" / "No mine was late!!" ✅
- "yep she did!" / "no not yet..." ✅
- “Ok cool!” / “Nah Wednesday does not work” ✅

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Groundedness - Grading Rules

2. **Replies that naturally extend the conversation are grounded.**

A reply may include some detail beyond what was explicitly stated, as long as it is plausible and flows naturally from the conversation without inventing specific facts (names, dates, places).

 - "I am not a fan." ✓
 - "i wish i could but that's a bit late 😭" ✓
3. **Rhetorical expressions and reactions are grounded.**

Exclamations, emojis, and conversational reactions that naturally fit the tone and moment of the conversation are grounded, even if they don't confirm specific content.

 - "Right?!" / "😂😂" ✓
 - "Lmaooo!!!" ✓
4. **Recipient nickname should be established in the input text or user profile**

 - The recipient nickname in a response would be considered **grounded** if you can identify it in the **input text** or **the user profile**.
 - If not, then the nickname would be considered hallucination – not grounded

Table of Contents

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Groundedness - Grading Rules

- .
- 5. **Follow-up questions are never grounded.**
A reply that primarily asks for more information is not grounded, regardless of how naturally it fits the conversation.
 - "Oh no, what happened?" ❌
 - "Have you tried calling them?" ❌
 - 6. **Specific fabric details not inferable from conversation are not grounded**
A reply that introduces specific facts like date, time, prices, or any other related details that are not clearly inferable from the context
 - Last message from PersonA: "Will you be here in 20 mins?"
Suggested response: "I will be there in 30 mins" ❌
 - How much are plates for an 08 avalanche by the way?
Suggested response: "I can check for you! For an 08, it shouldn't be too bad, probably around \$50-60 but I can look it up for sure when I'm back in" ❌

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Groundedness

	Input Text	Suggested Reply	Grade	Rationale
1	PersonA: "I'm heading to Bill's for lunch." Me: "Cool, I'll meet you there in 15." PersonA: "Sounds good. See you then."	"See you in 15."	 Grounded	The reply confirms information already stated in the conversation ("meet you there in 15"). It is a clear statement, not a question, and introduces no fabricated specific details.

Table of Contents

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Groundedness

	Input Text	Suggested Reply	Grade	Rationale
2	PersonA: The project deadline is Friday. I'm almost done with my part. Me: Great, I'll finalize the report by then. PersonA: "Perfect."	"Cool! "Deadline is Friday."	 Grounded	The reply confirms a key detail explicitly stated in the conversation (the Friday deadline). It is a factual statement, not a question, and introduces no fabricated specific details

Table of Contents

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Groundedness

	Input Text	Suggested Reply	Grade	Rationale
3	Me: do you play tennis? PersonA: i don't, why do you ask? Me: you look like someone who plays tennis PersonA: HAHAHAHA i'm literally loling PersonA: i don't know what someone who plays tennis looks like	"Wanna learn?"	✗ Not Grounded	"Wanna learn?" is primarily a follow-up question. It also suggests an action (learning tennis) that no participant expressed intent to pursue.

Table of Contents

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Groundedness

	Input Text	Suggested Reply	Grade	Rationale
4	Me: Hey Jake any news about the weekend? PersonA: Susan is coming Friday morning. But her mother asked if we can visit her that day, too. Me: Ok. Where is she now? PersonA: At home. She's fine but a bit bored... A nurse comes everyday but she can't go outside	"Let’s visit her."	 Grounded	The reply is a plausible response that naturally flows from the conversation.

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Groundedness

New Examples Added on May 6th

The below examples are new, and they mean to show you how to apply the updated Groundedness guidelines.

	Input Text	Suggested Reply	Grade	Rationale
5	username: she just know she needed a car. username: i can't keep being her chauffeur. PersonA: Oh damn. I forgot someone hit her car. PersonA: Did she get paid out?	"yep she did!" "no not yet..."	✔ Grounded	Both are plausible polar responses to a direct yes/no question in the conversation. "no not yet..." includes a small natural detail but no fabricated specifics. By contrast, a reply like "Yes, she got \$5,000 from the insurance company" would NOT be grounded — it invents specific financial details not mentioned.

Table of Contents

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Groundedness

	Input Text	Suggested Reply	Grade	Rationale
6	PersonA: Ok! username: actually I can start heading out now if ure ready. username: I can finish my work later 😇 PersonA: I can be there in 30!	"okii!!" "i wish i could but that's a bit late 😭"	✅ Grounded	Both replies naturally flow from the context. "okii!!" confirms readiness, and "i wish i could but that's a bit late 😭" expresses a constraint that, while not explicitly stated, is a plausible and natural response. No fabricated specifics are introduced. By contrast, a reply like "I'll meet you at the coffee shop on Main Street!" would NOT be grounded — it invents a specific location not mentioned.

Table of Contents

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Groundedness

	Input Text	Suggested Reply	Grade	Rationale
7	username: now this is the content we love to see. PersonA: Laughed at "Lmaooo why is my house a mess??!!! 🤨" PersonA: Laughed at "I hate pottery!!"	"Lmaooo!!!" "Right??!"	✅ Grounded	These are rhetorical expressions that naturally respond to the tone and moment of the conversation. They don't confirm specific facts but are clearly grounded in the shared reaction. By contrast, a reply like "I know, that video was posted by @username last Tuesday!" would NOT be grounded — it invents specific details not present in the conversation.

Table of Contents

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Groundedness

	Input Text	Suggested Reply	Grade	Rationale
8	_username_: Bro sites? PersonA: Oh yeah PersonA: Smol Frens sites PersonA: What are the Twitter and Instagram?	"Right??!" "I'll send em over in a few just double checking handles", "Let me check one sec think it's just smolfrens for both"	✗ Not Grounded	Both responses add specific fabricated details like "double checking handles" - which is a specific action stated for username, or "just smolfrens for both" - the context mentions "Smol frens sites" - but there is no evidence of them being instagram handles.
9	PersonA: I bought her a lions onsie _username_: She has one at my house PersonA: Yeah but isn't it 3-6 months PersonA: Go Lions!!!!	"No it's a newborn."	✗ Not Grounded	The response "No it's a newborn." is not grounded - as it assumes specific details (although related) that are not explicitly stated.

Table of Contents

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Groundedness

	Input Text	Suggested Reply	Grade	Rationale
10	PersonA: How many songs PersonA: How long comedy? _username_: 2 songs and your comedy song PersonA: That’s 30 minutes!	"No, 25 minutes."	✗ Not Grounded	"No, 25 minutes" is not grounded - as it responds with a specific duration that is not explicitly stated

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Contextual fit

Overview

The suggested reply should be contextually suitable for the conversation, addressing the main subject, intent, and theme appropriately.

Grading Question and Scales



Q. Is the suggested reply contextually suitable for the conversation?

- **Contextually Fit:** The suggested reply is suitable for the conversation, appropriately addressing the main subject, intent, and theme.
- **Partially Contextually Fit:** The suggested reply is suitable for part of the conversation — it addresses some messages, threads, or topics in the context but not all. Use this when the reply addresses an earlier message but not the most recent one, or when the reply fits the general theme but misses a specific element.
- **Contextually Misfit:** The suggested reply is unsuitable for the conversation, failing to address the main subject, intent, or theme, or introducing irrelevant content.

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Contextual fit



Callout 1: Contextual Fit Can Be Subjective

Contextual Fit can be subjective and challenging compared to the other dimensions. A reply's suitability often depends on social norms and the specific relationship between the speakers, which may be unknown to you.

For example, it's hard to know if a sarcastic or playful tone is appropriate if you can't determine the speakers' relationship.

Rule: When in Doubt, Be Cautious

If the relationship or context is unclear, always favor a reply that is conventionally appropriate (i.e., leaning toward formal or serious). Avoid rating highly on replies that depend on a specific, familiar relationship to be considered fitting.

- Example 4 further illustrates this callout.

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Contextual fit



Callout 2: Multi-Message Contexts

When a conversation contains multiple messages from Person A, or multiple distinct threads, a single reply may not address every message — and that is acceptable.

Apply **Partially Contextually Fit** when:

- The reply addresses an earlier message but ignores the most recent one, yet the earlier message is still a reasonable anchor for a reply.
- The reply fits the general conversational theme but omits one specific element.

Do not penalize a reply as Contextually Misfit solely because it does not address the very latest message, if it is still meaningfully engaged with the conversation.

Table of Contents

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Contextual fit

Examples

	Input Text	Suggested Reply	Grade	Rationale
1	Person A: I just got promoted to senior manager.	That's awesome! Congrats!	✅ Contextually Fit	The reply celebrates the achievement matching the context.
2	Person A: Can you send me the project files by EOD?	Yep, will do	✅ Contextually Fit	The reply directly addresses the request.
3	Person A: We gonna get a rental car Person A: Shit to crazy Person A: It was \$20 to go 5 mins PersonA: Y'all can ride with us lol	Solid! Omw!	✅ Contextually Fit	Solid Omw - would be a plausible response and we dont want to overfit based on the details like "future plan, not a current state"

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Contextual fit

Examples

	Input Text	Suggested Reply	Grade	Rationale
4	_username_: i told him to go lay down and go to sleep and night night _username_: now he plopped in front of my door Person A: awwwww Person A: Sarah is cute	"Deadlines aren't real lol"	✔ Contextually Fit	The participants might be talking about 2 different subjects so reply is plausible

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Contextual fit

Examples

	Input Text	Suggested Reply	Grade	Rationale
5	PersonA: You getting out to eat? username: Yea you here? PersonA: 2 min PersonA: Thousands of packages go down conveyer what's chances I see one for Duncan	Ok	🟡 Partially Contextually Fit	"Ok" might not address the most recent message or most recent topic shift (a seemingly humorous remark) shared by User. However, it does acknowledge the earlier time estimate provided by PersonA, making it a partial fit.
6	username: At Gina's? PersonA: Can you get a bowl w brown rice, no beans, chicken, cheese, pico, green sauce, sour cream, and love PersonA: Ur a gem PersonA: Yea I'm at Gina's	haha thanks love	🟡 Partially Contextually Fit	"Haha thanks love" is partially fit as it does not acknowledge the most recent confirmation given by PersonA which is the current discussion, but responds to the compliment "Ur a gem" - making it partially fit.

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Contextual fit

Examples

	Input Text	Suggested Reply	Grade	Rationale
7	PersonA: If your okay with tht username: Like what do you mean PersonA: Like maybe next thursday we do something me, Stacey, you, whoever is gonna be with you we go do something PersonA: Ugh im here and still been waiting	Yeah	🟡 Partially Contextually Fit	"Yeah" does not address the most recent message about PersonA's frustration with waiting which might warrant a reply. Instead, it responds to the earlier suggestion about making plans, where an agreement like "Yeah" would be appropriate - making it partially fit.

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Contextual fit

Examples

	Input Text	Suggested Reply	Grade	Rationale
8	Person A : My grandmother passed away yesterday	Oh no. Did you see the game?	✗ Contextually Misfit	The reply fails to respect the serious context.
9	Person A: I'm really stressed about the deadline	Deadlines aren't real lol	✗ Contextually Misfit	The relationship is unclear; a playful dismissal is not conventionally appropriate here. A more suitable reply would be: "Oh no, tell me about it."

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements


Step 4. Evaluate the generic quality of the reply.

Conciseness

Overview

The suggested reply needs to be complete and clear while remaining free of redundancy.

Grading Question and Scales



Q. Is the suggested reply concise?

- **Concise:** The suggested reply is complete, clear, and free of redundancy or unnecessary words.
- **Not Concise:** The suggested reply contains redundancy, wordiness, or unnecessary elaboration, or lacks completeness and clarity.

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Conciseness

Examples

	Input Text	Suggested Reply	Grade	Rationale
1	"Lunch at the cafe 11 pm tomorrow?"	"Sounds good!"	✓ Concise	The reply is clear and complete.
2	"Hey, are you free tonight? Want to go check out this movie?"	"Can't tonight, free tomorrow?"	✓ Concise	The reply is efficient and direct.
3	"How about adding the Museum of Modern Art to our itinerary?"	"Well, I think that maybe we could possibly consider looking at other options too"	✗ Not Concise	The reply is wordy and unclear.

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Conciseness

Examples

	Input Text	Suggested Reply	Grade	Rationale
4	"Can you make the meeting tomorrow?"	"I got your message about the meeting. I'm confirming I'll be there. The meeting is important."	✗ Not Concise	The reply sounds repetitive.
5	"Long time no see! Would you want to meet up sometime soon? Let me know your availability."	"Ok"	✗ Not Concise	The reply is too brief and vague. It doesn't seem complete.

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

Groundedness

Contextual Fit

Conciseness

Tone and Empathy

Naturalness

Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.


Tone & Empathy alignment

Overview

The suggested reply’s tone should appropriately reflect and respond to the recipient's expressed feelings and needs. Tone can be conveyed through:

- word choice
- punctuation
- formatting
- emojis
- length, and structure

Grading Question and Scales



Q. How well does the suggested reply's tone align with the emotional context of the conversation?

- **Aligned:** The suggested reply’s tone fully matches the emotional context and appropriately acknowledges the recipient's expressed feelings and needs.
- **Not Aligned:** The suggested reply’s tone slightly or completely mismatches the emotional context, failing to acknowledge or respond to the recipient's feelings.

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Tone & Empathy alignment



Callouts

- **Relationship and Formality:** When evaluating the Tone and Empathy dimension, consider the information provided by the Relationship and Formality labels.
- **If in doubt, be cautious:** If the context is unclear, always favor a reply that shows empathy and acknowledges the user's feelings.
 - Example 4 further illustrates this callout.

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Tone & Empathy alignment

Examples

	Input Text	Suggested Reply	Grade	Rationale
1	"I'm so nervous about the interview tomorrow" • Relationship: Roommate • Formality: Casual	"You got this! You're gonna crush it"	✅ Aligned	The reply acknowledge's the sender's feeling and shows encouragement.
2	"I just got engaged!!" • Relationship: Family • Formality: Casual	"OMG that's amazing!! So happy for you!"	✅ Aligned	The reply shows excitement and joy, which is well aligned with the context.

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Tone & Empathy alignment

Examples

	Input Text	Suggested Reply	Grade	Rationale
3	"I'm so frustrated with my boss right now" • Relationship: Professional • Formality: Casual	"Ugh, that's annoying. What happened?"	✅ Aligned	The reply acknowledges the sender's frustration and responds by asking what happened. Note: This example would be considered Not Grounded because the model is not expected to ask a follow-up question. In this scenario, a simple acknowledgement would be acceptable.
4	"I'm heartbroken. My partner broke up with me" • Relationship: Family • Formality: Neutral	"Breakups happen. Just move on"	❌ Not aligned	This is not the best response. The brevity of the reply can come across as cold and dismissive.

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Tone & Empathy alignment

Examples

	Input Text	Suggested Reply	Grade	Rationale
5	"I failed my driving test again" • Relationship: Friendship • Formality: Very informal	"Aw, you'll get it next time"	✗ Not aligned	This can be an edge case, given that one can imagine a user talks to a friend like this. On the other hand, one can also argue this reply sounds like a low-effort, slightly dismissive response. Given that we do not know for now the exact relationship between the sender and the recipient, please favor a response that acknowledges the user's frustration.

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Tone & Empathy alignment

Examples

	Input Text	Suggested Reply	Grade	Rationale
6	"Can you help me with my presentation?" • Relationship: Professional • Formality: Neutral	"Sure! I know you're presenting to the board tomorrow"	✅ Aligned	The reply shows tone and empathy alignment, although it makes up details that cannot be traced back to the input text.
7	"I'm going to the gym after work" • Relationship: Roommate • Formality: Casual	"Nice! I heard you're training for a marathon"	✅ Aligned	Remember the rule to grade each dimension independently. You'd penalize response hallucination under Groundedness .

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements


Step 4. Evaluate the generic quality of the reply.

Naturalness

Overview

The suggested reply should sound like something a real person would naturally say in this context, using authentic language and phrasing that is appropriate for the conversation's level of formality.

Grading Question and Scales



Q. How natural does the suggested reply sound?

- **Natural:** The suggested reply sounds like authentic human conversation, using natural language, phrasing, and expressions appropriate for the context and formality level.
- **Unnatural:** The suggested reply sounds awkward, unnatural, or overly formulaic—with forced phrasing or mechanical language that doesn't resemble authentic human communication.

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Naturalness

Examples

	Input Text	Suggested Reply	Grade	Rationale
1	"I've been stressed about this upcoming big travel plan."	"Yeah same, I'd be stressed too"	✔ Natural	The reply uses conversational, authentic phrasing.
2	"See, I told you, my brother looks like Tom Cruise."	"Haha no way, that's wild"	✔ Natural	The reply sounds casual and conversational.
3	"Because of that, I've been really worried these days."	"I acknowledge your aforementioned concern"	✘ Unnatural	The reply sounds overly formal for casual messaging. Improved version: "Yeah, I get why you're worried about that"

- Introduction
- Workflow
 - Step 1 - Review Input
 - Step 2 - Categorize Input
 - Step 3 - Harmfulness
 - Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
 - Step 5 - Personalization
 - Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Naturalness

Examples

	Input Text	Suggested Reply	Grade	Rationale
4	"That BBQ on the other side of the city is soooo good!"	"The gustatory experience was satisfactory"	✗ Unnatural	The reply sounds overly formal for casual messaging. Improved version: "Yup, the food was pretty good!"
5	"Hey, I've just shared the doc."	"PROCESSING. STANDBY FOR RESPONSE."	✗ Unnatural	The reply looks like it's machine-generated.

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Localization

Overview

The suggested reply should feel like it was written by a native speaker in the target region. It adapts to local language, culture, and context.

A response fails the localization dimension if it feels "foreign," even if the grammar is technically correct.

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization


Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Localization

Grading Question and Scales



Q. Does the response display content that is appropriate and relevant for your language and region?

- **Yes:** The response is fully appropriate for my language and region
- **No:** The response contains something inappropriate or irrelevant for my language and region
 - Awkward or clearly machine-translated phrasing
 - Incorrect mixed languages in the response
 - Wrong date format (e.g., MM/DD vs DD/MM), units of measurement, or number formatting
 - Wrong punctuation format for my language/region
 - Scrambled symbols (, □, ???, ?)
 - Content, examples, or references that don't fit my culture/region

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Localization

Instructions

-  If **all** of the following are true, select Yes for the Localization question:
- Input text is in my language with no obviously incorrect or untranslated portions
 - Dates, numbers and units of measurement display in my local format
 - No garbled characters appear
 - All content, including cultural references, punctuations, and examples, make sense for my region

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 4. Evaluate the generic quality of the reply.

Localization

Instructions




Guideline Callouts

1. Below scenarios are expected and NOT Localization issues:

- English proper names of companies, products, games, movies, etc **will appear frequently in other languages.**
- Some languages such as Hinglish use mixed languages.
- It’s possible that some scenarios use the base language structurally but include a small untranslated portion.
 - For example, “Send 你好嗎 to Ching Liu.”

2. Examples of wrong punctuation format

- ✗ Cómo estás? Qué sorpresa! **Vs** ✔ ¿Cómo estás? ¡Qué sorpresa!
- ✗ Il a dit "bonjour" en souriant. **Vs** ✔ Il a dit « bonjour » en souriant.
- ✗ Powiedział "to dobrze" i wyszedł. **Vs** ✔ Powiedział „to dobrze" i wyszedł.



Steps and Requirements - **Step 5**

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 5. Evaluate the personalization quality of the reply.

Overview

You have evaluated the generic quality of the suggested reply. Now, you'll be evaluating the personalization quality. There is only one dimension to evaluate in this step.

Dimension	Description
Personalization	Personalization measures how well the suggested reply captures the user's unique communication style while considering the conversation context.

The User Profile

The suggested reply should *broadly* capture the user's communication style (also known as **User Profile**) while still considering the conversation context.

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 5. Evaluate the personalization quality of the reply.

The **User Profile** will include the components below, which will be displayed for you in the grading tool.

	How to understand the component?
Common phrases	Phrases or messages (sometimes with emojis) commonly used by the users.
Abbreviations and Contractions	Abbreviations: Shortened forms of words or phrases commonly used by the user. Contractions: The user's preference to use contraction over the expanded form (such as preferring "cant" over "cannot", "lets" over "let us"). Note: The use of "can't" or "let's" is acceptable and would still be considered personalized.
Punctuations	The user's preferred style and frequency of punctuation marks.
Nickname	Nicknames the user uses for the participant(s). This field can be empty sometimes, meaning there's not nickname info to be evaluated.

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

Groundedness

Contextual Fit

Conciseness

Tone and Empathy

Naturalness

Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 5. Evaluate the personalization quality of the reply.

A sample user profile can look like the below:

User Profile	
Common phrases	Punctuations
Acknowledgment (neutral): ['Hmm', 'I know', 'Oh, like that', 'It does', 'I hear ya.', 'Oh I didn't know that', 'No idea.', 'Ya ok']	Capitalization: Rare: ["Sentence-Start Capitalization","Emphasis","Acronyms & Initialisms"]
Agreement & confirmation: ['I know', 'It does', 'Yeah I know the feeling', 'I hear ya.', 'Tonight we will be.', "Yeah it's on there", 'Ya ok', 'And I amm']	Periods: Rare: ["Sentence Termination","Abbreviation Marking"]
Acknowledging Acknowledgements: ['I know', 'I hear ya.', 'It does', 'Hmm', 'Ya ok']	Exclamation Points: Frequent: ["Excitement/Emphasis","Intensified Exclamation (multiple exclamation marks)"] Rare: ["Stylistic Habit"]
Disagreement & Rejection: ["'Whatever'", "'Idc'", "'I'm moving on'", 'Was that necessary', 'I don't think it'd be the end of the world', 'But I know that prob wouldn't go down well.', 'Nothing really', 'Not that they will', 'AirPods just look ridiculous', 'Never heard of lite coin', 'No one on this crew has heard of it', 'Probably just a fad.', 'Ya ok, but I only really agree with like one of those ideas']	Question Marks: Frequent: ["Direct Questions"] Rare: ["Intensity (repeated question marks)"]
Politeness: ['Is there anything you want?', 'Haha, I'm just being dumb', 'Maybe mentos can do you a favor for a change', '🙄']	Ellipses: Frequent: ["Trailing or Pausing"] Rare: ["Omission"]
Expressing emotion: ['😞', 'Haha, I'm just being dumb', 'I hear ya.', 'has kinda screwed me.', 'AirPods just look ridiculous', '😞', 'Omg already doing my fucking head in', 'Prob bc I'm tired', 'Yeah I know the feeling', 'Ridiculously silly drunk.']	Commas: Rare: ["Structural Use","Stylistic Pauses"]
Greeting & closing: ['👋']	
Abbreviations, Contractions	
Abbreviations: ['atm', 'idc', 'omg']	
Contractions: ['The user has avoids contractions and generally prefers full words.']	

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 5. Evaluate the personalization quality of the reply.

Grading Question and Scales



Q. How well does the suggested reply reflect the user's communication style?

- **Personalized:** The reply clearly reflects the user's communication style, tone, and language patterns.
- **Contextually Adapted:** The reply deviates from the user's typical style *but adapts appropriately* to the conversation's tone, formality, or context, as the user would naturally do.
- **Generic:** The reply is generic, lacking the user's unique style, tone, or other communication characteristics typically used. A personalized reply **would have been appropriate** based on past user responses.
- **Mismatched:** The reply directly contradicts the user's established communication style, such as their tone, punctuation usage, or fails to use context-relevant phrases.

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 5. Evaluate the personalization quality of the reply.



More context for analysts to align with the engineering team:

For the engineering team:

- **Personalized** and **Contextually Adapted**: This receives a similar score because the assistant should not be penalized for providing a relevant, context-appropriate reply.
- **Generic**: This incurs a penalty when a reply could have been personalized but was not.

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 5. Evaluate the personalization quality of the reply.

How to evaluate the Personalization dimension?

It is acknowledged by the engineering team that this dimension can be subjective and context-based. To guide your grading, we have added the grading principles below:

 **Personalization Dimension - Grading Principles**

- **Guidance, Not Law:** Consider the guidelines for this dimension as a framework to inform your decisions, not strict mandates. Human communication is fluid, and our goal is to capture its nuances.
- **User Profile as a Tendency:** The user profile (capitalization, punctuation, emoji use, abbreviations, etc..) describes what a user *often* does, not what they *always* do. It's "including but not limited to" these elements. Look for patterns, but recognize exceptions.
- **Context is Key:** The relevance of specific profile components (e.g., emoji vs. capitalization) is dynamic and depends entirely on the conversation's context. Use your common sense and the provided examples to weigh which components are most significant for the task at hand.
- **Leverage Your Insight:** We value your ability to extract patterns and make reasonable, informed decisions. Think beyond surface-level checks and consider the overall "feel" of the reply in relation to the user's typical communication style.
- **Aim for Likelihood:** Your decision is about assessing the likelihood that a reply aligns with the user's personalized style. A "good" personalization might not perfectly match every single aspect, but it should feel authentic and appropriate.

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 5. Evaluate the personalization quality of the reply.

Examples

User Profile

Capitalization	▪ frequent: none ▪ rare: Sentence-Start Capitalization, Emphasis
Abbreviations	▪ "u", "omg", "lol", "ur", "btw"
Punctuations	▪ Question marks: direct questions (frequent), intensity (rare); expressive questions (occasional) ▪ Ellipses: trailing or pausing (frequent); omission (rare) ▪ Exclamation points: excitement, emphasis, intensified (frequent); stylistic habit (rare)

Table of Contents

- Introduction
- Workflow
- Step 1 - Review Input
- Step 2 - Categorize Input
- Step 3 - Harmfulness
- Step 4 - Generic Qualities
 - Groundedness
 - Contextual Fit
 - Conciseness
 - Tone and Empathy
 - Naturalness
 - Localization
- Step 5 - Personalization
- Step 6 - Comparison

Steps and Requirements

Step 5. Evaluate the personalization quality of the reply.

Examples

User Profile

Contraction	high preference of using: "cant", "well", "lets"
Common phrases	Acknowledgment: "lol knew u'd love it", "yes", "true", "yup", "lol good", "will do", "knew u'd like it" Agreement & confirmation: "omg deal", "duh", "yes", "true", "omg yes!!!", "yup", "perfect combo", "always", "will do" Politeness: "anytime ❤️", "for u ❤️", "always 🥰🥰🥰", "all for u ❤️", "want me to bring some later?", "hehe always for u ❤️" Greeting & closing: "hey babe!", "morning 🥰", "hey", "hey love 🥰", "haha morning babe ❤️", "night babe 🥰" Expressing emotion: "omg yes!!!", "omg 😍", "missed u today 😭", "🥰❤️", "can't wait 😍", "thinking of u 🥰", "love u more ❤️", "hehe 🥰", "crying lol 🤔", "traffic sucks 😞"
Nickname	bae, babe

Steps and Requirements

Step 5. Evaluate the personalization quality of the reply.

Input Text / Chat Context	Suggested Reply	Grade	Rationale
"See, I told you, my brother looks like Tom Cruise."	"Haha no way, that's wild!"	Personalized	
	"Haha no way, that's wild"	Personalized	This can be an edge case, given that based on the user profile, the user tends to use exclamation marks or emojis to express emotions. However, one can also argue that personalization here comes from the fact that the user use words like "that's wild," which could reflect an informal/casual style. The engineering team would consider an example like this as Personalized.
	"Haha no way, that's wild 🤔 "	Personalized	
	"Haha no way, that's wild" 😎	Personalized	Although 😎 doesn't appear in the profile, it's possible that a user uses this emoji in this chat context (to express "wow, cool" or "that's exciting!")

Steps and Requirements

Step 5. Evaluate the personalization quality of the reply.

Input Text / Chat Context	Suggested Reply	Grade	Rationale
"I'm so nervous about the interview tomorrow"	"You got this! I'm here for you."	Contextually Adapted	It can be hard to tell the exact level of formality of the chat context in this example. The tone in the input seems neutral and not overly informal (unlike "I'm soooo nervous about the interview tomorrow!"). In this context, it is appropriate for the user to match the conversation's tone, formality, or context in a way the user would naturally adjust.
	"You got this! I'm here for u ❤️"	Personalized	
	"hey, u got this. all for u ❤️"	Personalized	
	"I understand that interviews can be daunting, and I wish you the best of luck tomorrow."	Mismatched	The sentence sounds elaborate and formal, mismatching the style suggested in the user profile.

Table of Contents

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

Groundedness

Contextual Fit

Conciseness

Tone and Empathy

Naturalness

Localization


Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 5. Evaluate the personalization quality of the reply.

Input Text / Chat Context	Suggested Reply	Grade	Rationale
"Hey, are you free tonight? Want to go check out this movie?"	"Cant tonight, free tomorrow?"	Personalized	The User Profile suggests a high preference for "cant," which is used to personalize this reply
	"Can't tonight, free tomorrow?"	Personalized	The User Profile says high preference of using: "cant", "well", "lets." Note that this requirement means the user prefers contraction "cant" over "cannot." So either "can't or "cant" is acceptable and considered personalized from the engineering perspective.
	"yup ❤️"	Personalized	
	YES! COUNT ME IN!	Mismatched	The user profile doesn't suggest a tendency to use all caps in a reply.
	OK	Mismatched	The user profile suggests a style that's expressive so a reply like "OK" seems too brief and cold.
	"PROCESSING. STANDBY FOR RESPONSE."	Mismatched	This appears to be a machine-generated reply, not personalized at all.



Steps and Requirements - **Step 6**

Introduction

Workflow

Step 1 - Review Input

Step 2 - Categorize Input

Step 3 - Harmfulness

Step 4 - Generic Qualities

 Groundedness

 Contextual Fit

 Conciseness

 Tone and Empathy

 Naturalness

 Localization

Step 5 - Personalization

Step 6 - Comparison

Steps and Requirements

Step 6. Compare Responses in Pair

In Step 6, you are comparing the **collective quality** of the suggestions in **Assistant A** and **Assistant B**.


Think of it this way:

- Assistant A and Assistant B are two different sources of suggestions.
- Your task is to decide which option, as a whole, gave you a better set of suggestions.
- Example: If Assistant A has one suggestion and Assistant B has two, you are comparing A1 vs. the combined group of (B1 + B2).

How to Compare

- **Group the responses.** Look at all of Assistant A's suggestions (A1, A2) as one complete group. Do the same forAssistant B (B1, B2).
- **Compare the groups.** Ask yourself: "Overall, which group of suggestions is more helpful and useful? The group from A or the group from B?"

Share any feedback you'd like the engineering team to be aware of and submit the task.



Congratulations

You have completed the guidelines for Writing Tools –
Personalized Smart Replies!