

WT – Mail Smart Reply Feedback

This feedback is based on Inter-Analyst Agreement (IAA). That is, how much do different people analyzing the same task agree with each other? The observations and feedback below are sorted based on the most recurring issues, i.e. from the lowest agreement scores.

1. Personalization (HIGH RISK)

This is the most widespread and severe issue in the data. You should revisit how you're applying the Personalization dimension.

The guidelines are explicit that a Personalized response must reflect the user's writing style patterns (vocabulary, formatting, padding, sentence structure), their Mail Component Patterns (salutations, openings, closings, sign-offs, signature), and the relationship context. A response that is functionally correct but uses a generic salutation, wrong sign-off, or misses the user's typical formatting should be rated Generic, not Personalized.

Key reminders from the guidelines:

- The user profile describes tendencies, not strict rules, but you should look for meaningful alignment, not just surface-level checks.
- If a user profile shows "strong_formatting" and the draft uses no lists or structure when the prompt explicitly asked for one, that is Generic.
- If the draft uses a completely different sign-off than what the profile shows (e.g., "Thanks" vs. "Thank You"), that also pushes toward Generic.
- You should think about the overall "feel"; does this read as authentically from this particular person, or could anyone have written it?

Calibrate against the given examples in Step 6, especially the contrast between Reply 1 (Personalized) and Reply 2 (Generic) in the Marija example, which is the clearest illustration of what "missing style elements" looks like in practice, and what would be the more common scenarios across tasks.

2. Groundedness (HIGH RISK)

The v1.3 update introduced a three-tier scale - Grounded, Partially Grounded, Not Grounded. The "Partially Grounded" tier is new; make sure you're clear on the correct approach applicable to the in-task scenarios.

The critical distinctions to reinforce:

- Standard professional phrases ("Thanks for checking in," "Hope you are doing well," pleasantries, conventional closings) are **Grounded** - do not penalize these.
- A draft that gets the core message right but adds one minor invented detail (a specific time not in the prompt, a fabricated duration, a reference to a prior conversation that didn't happen) is **Partially Grounded**, not Not Grounded.
- A draft that introduces an invented fact that changes the meaning of the message, or brings in a topic entirely unrelated to the instruction, is **Not Grounded**.
- Rule 5 is particularly important: if the recipient's name cannot be found anywhere in the input (prompt, prior email, user profile, additional info), the response should be rated Partially Grounded, not Not Grounded — as long as the core message is correct.
- As long as a name is found in the task (user prompt, previous email, additional info or profile), it would be marked as **Grounded**, even if contradicting the user instruction. For example, prompt

says “Email Mark...”, but the profile says “John” and the response starts with “Hi John”. It’s **Grounded** but **Contextually Misfit**.

Reserve **Not Grounded only for clear failures**. Minor invented details with the core message intact = Partially Grounded.

3. Instruction Adherence & Contextual Fit (HIGH RISK)

This dimension also now has a three-tier scale, and the same calibration risk applies. You should be careful not to rate a response as Not Followed or Misfit when Partially Followed and Fit is more accurate. The key rules:

- If a draft fulfills the primary instruction but omits a secondary element, or responds to most but not all key points in the prior email, or has a minor formality mismatch — that is Partially Followed and Fit.
- Not Followed or Misfit is reserved for cases where a critical element is missing, the instruction is contradicted, or the formality gap is significant (e.g., casual tone in a clearly formal professional thread).

You should also remember: when there is a prior email, you need to evaluate **both** instruction adherence and contextual fit together. A response can follow the instruction but fail to address the prior email's key points. That would be Partially Followed and Fit at best.

4. Locale-Specific Issues

For **zh-CN** and **zh-HK** specifically: the Localization dimension is HIGH RISK, which is unique to these locales. You should pay close attention to script-level issues, wrong punctuation format or formatting that does not fit Chinese-language conventions. Do NOT apply LE cert localization rules. Writing Tools have their own localization rules outlined in the guides. Grammar/Spelling do not fall under localization. Most other locales are at MEDIUM RISK, so please observe the exact expectations on Localization and be cautious in your evaluations.

For **vi-VN**: Naturalness is also flagging as HIGH RISK, which is unusual compared to other locales. This may reflect model output that is grammatically acceptable but reads as translated or stilted in Vietnamese. You should apply the Naturalness standard carefully, the question is whether the draft sounds like a real person wrote it in that language, not just whether the grammar is correct.

5. What's Working - Don't Overcorrect

Harmful and Tone are consistently LOW RISK across all locales. You should not feel pressure to find issues in these dimensions where none exist. Tone in particular should only be rated Not Aligned when there is a clear mismatch — a cold, robotic reply to someone expressing frustration, or a casual tone in a formal professional context. If the tone is broadly appropriate, it should pass.
